

RHL LAB ALL SOLUTIONS

Lab-1: Improving Command-line Productivity

1. Create the /home/student/bin/bash-lab script file on workstation.
 1. On workstation, create the /home/student/bin/ folder if needed.

```
[student@workstation ~]$ mkdir -p /home/student/bin
```

2. Use **vim** to create and edit the /home/student/bin/bash-lab script file.

```
[student@workstation ~]$ vim ~/bin/bash-lab
```

3. Insert the following text and save the file.

```
#!/bin/bash
```

4. Make your script file executable.

```
[student@workstation ~]$ chmod a+x ~/bin/bash-lab
```

You can use **sudo** without requiring a password on the servera and serverb hosts. Remember to use a loop to simplify your script. You can also use multiple **grep** commands concatenated with the use of the pipe character (|).

1. Use **vim** to open and edit the /home/student/bin/bash-lab script file.

```
[student@workstation ~]$ vim ~/bin/bash-lab
```

2. Append the following lines in bold to the /home/student/bin/bash-lab script file.

Note

The following is an example of how you can achieve the requested script. In Bash scripting, you can take different approaches and obtain the same result.

```
#!/bin/bash  
  
#  
USR='student'  
OUT='/home/student/output'  
  
#
```

```

for SRV in servera serverb
do
ssh ${USR}@${SRV} "hostname -f" > ${OUT}-${SRV}
echo "#####" >> ${OUT}-${SRV}
ssh ${USR}@${SRV} "lscpu | grep '^CPU'" >> ${OUT}-${SRV}
echo "#####" >> ${OUT}-${SRV}
ssh ${USR}@${SRV} "grep -v '^$' /etc/selinux/config|grep -v '^#'" >>
${OUT}-${SRV}
echo "#####" >> ${OUT}-${SRV}
ssh ${USR}@${SRV} "sudo grep 'Failed password' /var/log/secure" >>
${OUT}-${SRV}
echo "#####" >> ${OUT}-${SRV}
done

```

1. Execute the **/home/student/bin/bash-lab** script, and review the output content on workstation.
 1. On workstation, execute the **/home/student/bin/bash-lab** script.

```
[student@workstation ~]$ bash-lab
```

2. Review the content
of **/home/student/output-servera** and **/home/student/output-serverb**.

```

3. [student@workstation ~]$ cat /home/student/output-servera
4. servera.lab.example.com
5. #####
6. CPU op-mode(s):      32-bit, 64-bit
7. CPU(s):              2
8. CPU family:          21
9. CPU MHz:             2294.670
10. #####
11. SELINUX=enforcing
12. SELINUXTYPE=targeted
13. #####
14. Mar 21 22:30:28 servera sshd[3939]: Failed password for invalid user
operator1 from 172.25.250.9 port 58382 ssh2
15. Mar 21 22:30:31 servera sshd[3951]: Failed password for invalid user
sysadmin1 from 172.25.250.9 port 58384 ssh2
16. Mar 21 22:30:34 servera sshd[3953]: Failed password for invalid user
manager1 from 172.25.250.9 port 58386 ssh2

```

```
#####
[student@workstation ~]$ cat /home/student/output-serverb
serverb.lab.example.com
#####
CPU op-mode(s):      32-bit, 64-bit
CPU(s):              2
CPU family:          6
CPU MHz:              2294.664
#####
SELINUX=enforcing
SELINUXTYPE=targeted
#####
Mar 21 22:30:37 serverb sshd[3883]: Failed password for invalid user
operator1 from 172.25.250.9 port 39008 ssh2
Mar 21 22:30:39 serverb sshd[3891]: Failed password for invalid user
sysadmin1 from 172.25.250.9 port 39010 ssh2
Mar 21 22:30:43 serverb sshd[3893]: Failed password for invalid user
manager1 from 172.25.250.9 port 39012 ssh2
#####
```

Lab2: Tuning System Performance

1. Change the current tuning profile for serverb to balanced, a general non-specialized tuned profile.
 1. From workstation, open an SSH session to serverb as student user. The systems are configured to use SSH keys for authentication, so a password is not required.

```
2. [student@workstation ~]$ ssh student@serverb
3. ...output omitted...
4. [student@serverb ~]$
```

5. Use **yum** to confirm that the tuned package is installed.

```
6. [student@serverb ~]$ yum list tuned
7. ...output omitted...
8. Installed Packages
9. tuned.noarch                2.10.0-15.el8
   @anaconda
```

10. Use the **systemctl is-active tuned** command to display the tuned service state.

```
11. [student@serverb ~]$ systemctl is-active tuned
12. active
```

13. List all available tuning profiles and their descriptions. Note that the current active profile is virtual-guest.

```
14. [student@serverb ~]$ sudo tuned-adm list
15. [sudo] password for student: student
16. Available profiles:
17. - balanced                - General non-specialized tuned profile
18. - desktop                 - Optimize for the desktop use-case
19. - latency-performance    - Optimize for deterministic performance at
    the cost of
20.                          increased power consumption
21. - network-latency        - Optimize for deterministic performance at
    the cost of
22.                          increased power consumption, focused on
    low latency
23.                          network performance
24. - network-throughput     - Optimize for streaming network throughput,
    generally
25.                          only necessary on older CPUs or 40G+
    networks
26. - powersave             - Optimize for low power consumption
27. - throughput-performance - Broadly applicable tuning that provides
    excellent
28.                          performance across a variety of common
    server workloads
29. - virtual-guest          - Optimize for running inside a virtual
    guest
30. - virtual-host           - Optimize for running KVM guests
31. Current active profile: virtual-guest
```

32. Change the current active tuning profile to the balanced profile.

```
[student@serverb ~]$ sudo tuned-adm profile balanced
```

33. List summary information of the current active tuned profile.

Use the **tuned-adm profile_info** command to confirm that the active profile is the balanced profile.

```
[student@serverb ~]$ sudo tuned-adm profile_info
```

Profile name:

balanced

Profile summary:

General non-specialized tuned profile

...output omitted...

1. Two processes on serverb are consuming a high percentage of CPU usage. Adjust each process's nice level to 10 to allow more CPU time for other processes.

1. Determine the top two CPU consumers on serverb. The top CPU consumers are listed last in the command output. CPU percentage values will vary.

```
2. [student@serverb ~]$ ps aux --sort=pcpu
3. USER          PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME
   COMMAND
4. ...output omitted...
5. root           2983  100  0.0 228360 1744 ?        R<    21:08   0:23
   md5sum /dev/zero
6. root           2967  101  0.0 228360 1732 ?        RN    21:08   0:23
   sha1sum /dev/zero
7. [student@serverb ~]$
```

8. Identify the current nice level for each of the top two CPU consumers.

```
9. [student@serverb ~]$ ps -o pid,pcpu,nice,comm \
10. $(pgrep sha1sum;pgrep md5sum)
11.  PID %CPU  NI  COMMAND
12. 2967 99.6   2  sha1sum
13. 2983 99.7  -2  md5sum
```

14. Use the **sudo renice -n 10 2967 2983** command to adjust the nice level for each process to 10. Use PID values identified in the previous command output.

```
15. [student@serverb ~]$ sudo renice -n 10 2967 2983
16. [sudo] password for student: student
17. 2967 (process ID) old priority 2, new priority 10
18. 2983 (process ID) old priority -2, new priority 10
```

19. Verify that the current nice level for each process is 10.

```
20. [student@serverb ~]$ ps -o pid,pcpu,nice,comm \
21. $(pgrep sha1sum;pgrep md5sum)
```

```
22.  PID %CPU  NI  COMMAND
23.  2967  99.6  10  sha1sum
24.  2983  99.7  10  md5sum
```

25. Exit from serverb.

```
26. [student@serverb ~]$ exit
27. logout
28. Connection to serverb closed.
```

```
[student@workstation ~]$
```

Lab3: Controlling Access to Files with ACLs

1. The cases directory and its contents should belong to the managers group. New files added to the cases directory should automatically belong to the managers group. The user and group owners for the existing files should have read and write permission, and other users should have no permission at all.

Note

Hint: Do not use **setfacl**.

1. Log in to serverb as the student user.

```
2. [student@workstation ~]$ ssh student@serverb
3. ...output omitted...
```

```
[student@serverb ~]$
```

4. Use the **sudo -i** command to switch to the root user. The password for the student user is student.

```
5. [student@serverb ~]$ sudo -i
6. [sudo] password for student: student
```

```
[root@serverb ~]# Flab
```

7. Use the **chgrp** command to recursively update group ownership on the directory and its contents.

```
[root@serverb ~]# chgrp -R managers /shares/cases
```

8. Use the **chmod** command to update the set-GID flag on the directory.

```
[root@serverb ~]# chmod g+s /shares/cases
```

9. Use **chmod** to update all existing file permissions to `rw` for owner and group.

```
[root@serverb ~]# chmod 660 /shares/cases/*
```

1. Add ACL entries to the `cases` directory (and its contents) that allow members of the `contractors` group to have read/write access on the files and execute permission on the directory. Restrict the `contractor3` user to read access on the files and execute permission on the directory.

1. Use **setfacl** to recursively update the existing `cases` directory and its contents. Grant the `contractors` group read, write, and conditional execute permissions.

```
[root@serverb ~]# setfacl -Rm g:contractors:rwX /shares/cases
```

2. Use **setfacl** to recursively update the existing `cases` directory and its contents. Grant the `contractor3` user read and conditional execute permissions.

```
[root@serverb ~]# setfacl -Rm u:contractor3:rX /shares/cases
```

1. Add ACL entries that ensure any new files or directories in the `cases` directory have the correct permissions applied for *all* authorized users and groups.

1. Use **setfacl** to update the *default* permissions for members of the `contractors` group. Default permissions are read, write, and execute (needed for proper subdirectory creation and access).

```
[root@serverb ~]# setfacl -m d:g:contractors:rwX /shares/cases
```

2. Use **setfacl** to update the *default* permissions for the `contractor3` user. Default permissions are read and execute (needed for proper subdirectory access).

```
[root@serverb ~]# setfacl -m d:u:contractor3:rX /shares/cases
```

1. As the `root` user, use **ls** to check the `cases` directory and its content. Look for group ownership, directory and file permissions. The `"s"` in the group file permissions indicates the set-GID flag is set, and the `"+"` indicates that ACL entries exist. At the end exit from the `root` user session.

```
2. [root@serverb ~]# ls -ld /shares/cases
```

```
3. drwxrws---+ 2 root managers 46 Mar 29 00:40 /shares/cases
```

```
4. [root@serverb ~]# ls -l /shares/cases
```

```
5. total 8
6. -rw-rw----+ 1 root managers 44 Mar 29 00:33 backlog.txt
```

```
-rw-rw----+ 1 root managers 46 Mar 29 00:33 shortlist.txt
```

7. Use **getfacl** and review its output. Look for the named user and named group entries in both the standard and default ACL.

```
8. [root@serverb ~]# getfacl /shares/cases
9. # file: shares/cases
10.# owner: root
11.# group: managers
12.# flags: -s-
13.user::rwx
14.user:contractor3:r-x
15.group::rwx
16.group:contractors:rwx
17.mask::rwx
18.other:---
19.default:user::rwx
20.default:user:contractor3:r-x
21.default:group::rwx
22.default:group:contractors:rwx
23.default:mask::rwx
24.default:other:---
25.
26. [root@serverb ~]# exit
```

```
logout
```

27. Switch to the `manager1` user and perform the following operations. Check that you get the expected access behavior.

```
28. [student@serverb ~]$ su - manager1
29. Password: redhat
30. [manager1@serverb ~]$ cd /shares/cases
31. [manager1@serverb cases]$ echo hello > manager1.txt
32. [manager1@serverb cases]$ cat shortlist.txt
33. ###Shortlist of Clients to call###TEMPLATE###
34. [manager1@serverb cases]$ mkdir manager1.dir
```



```
35. [manager1@serverb cases]$ echo hello > manager1.dir/test.txt
36. [manager1@serverb cases]$ ls -ld manager1.dir
37. drwxrws---+ 2 manager1 managers 22 Mar 29 00:59 manager1.dir
38. [manager1@serverb cases]$ ls -l manager1.dir
39. total 4
40. -rw-rw----+ 1 manager1 managers 6 Mar 29 00:59 test.txt
41. [manager1@serverb cases]$ getfacl manager1.dir
42. # file: manager1.dir/
43. # owner: manager1
44. # group: managers
45. # flags: -s-
46. user::rwx
47. user:contractor3:r-x
48. group::rwx
49. group:contractors:rwx
50. mask::rwx
51. other:---
52. default:user::rwx
53. default:user:contractor3:r-x
54. default:group::rwx
55. default:group:contractors:rwx
56. default:mask::rwx
57. default:other:---
58.
59. [manager1@serverb cases]$ exit
```

logout

60. Switch to the contractor1 user and perform the following operations. Check that you get the expected access behavior.

```
61. [student@serverb ~]$ su - contractor1
62. Password: redhat
63. [contractor1@serverb ~]$ cd /shares/cases
64. [contractor1@serverb cases]$ echo hello > manager1.txt
65. [contractor1@serverb cases]$ cat shortlist.txt
66. ###Shortlist of Clients to call###TEMPLATE###
67. [contractor1@serverb cases]$ mkdir contractor1.dir
```

```
68. [contractor1@serverb cases]$ echo hello > contractor1.dir/test.txt
69. [contractor1@serverb cases]$ ls -ld contractor1.dir
70. drwxrws---+ 2 contractor1 managers 22 Mar 29 01:05 contractor1.dir
71. [contractor1@serverb cases]$ ls -l contractor1.dir
72. total 4
73. -rw-rw----+ 1 contractor1 managers 6 Mar 29 01:07 test.txt
74. [manager1@serverb cases]$ getfacl contractor1.dir
75. # file: contractor1.dir/
76. # owner: contractor1
77. # group: managers
78. # flags: -s-
79. user::rwx
80. user:contractor3:r-x
81. group::rwx
82. group:contractors:rwx
83. mask::rwx
84. other:---
85. default:user::rwx
86. default:user:contractor3:r-x
87. default:group::rwx
88. default:group:contractors:rwx
89. default:mask::rwx
90. default:other:---
91.
92. [contractor1@serverb cases]$ exit
```

logout

93. Switch to the contractor3 user, and perform the following operations. Check that you get the expected access behavior.

```
94. [student@serverb ~]# su - contractor3
95. Password: redhat
96. [contractor3@serverb ~]# cd /shares/cases
97. [contractor3@serverb cases]# echo hello > contractor3.txt
98. -bash: contractor3.txt: Permission denied
99. [contractor3@serverb cases]# cat shortlist.txt
100.  ###Shortlist of Clients to call###TEMPLATE###
```

```
101. [contractor3@serverb cases]# mkdir contractor3.dir
102. mkdir: cannot create directory 'contractor3.dir': Permission denied
103. [contractor3@serverb cases]# cat manager1.dir/test.txt
104. hello
105. [contractor3@serverb cases]# cat contractor1.dir/test.txt
106. hello
107. [contractor3@serverb cases]# exit
108. logout
```

```
[student@serverb ~]#
```

109. Log off from serverb

```
110. [student@serverb ~]# exit
111. logout
112. Connection to serverb closed.
```

```
[student@workstation ~]$
```

Lab4: Managing SELinux Security

1. Log in to serverb as the root user.
 1. Use the ssh command to log in to serverb as the student user. The systems are configured to use SSH keys for authentication, so a password is not required.

```
2. [student@workstation ~]$ ssh student@serverb
3. ...output omitted...
```

```
[student@serverb ~]$
```

4. Use the sudo -i command to switch to the root user. The password for the student user is student.

```
5. [student@serverb ~]$ sudo -i
6. [sudo] password for student: student
```

```
[root@serverb ~]#
```

1. Launch a web browser on workstation and browse to <http://serverb/lab.html>. You will see the error message: You do not have permission to access /lab.html on this server.
- 2.

3. Research and identify the SELinux issue that is preventing Apache from serving web content.

1. Using the less command, view the contents of /var/log/messages. Use the / key and search for sealert. Use the q key to quit the less command.

```
2. [root@serverb ~]# less /var/log/messages
```

3. Mar 28 10:19:51 serverb setroubleshoot[27387]: SELinux is preventing /usr/sbin/httpd from getattr access on the file /lab-content/lab.html. For complete SELinux messages run: sealert -l 8824e73d-3ab0-4caf-8258-86e8792fee2d

```
Mar 28 10:19:51 serverb platform-python[27387]: SELinux is preventing /usr/sbin/httpd from getattr access on the file /lab-content/lab.html.#012#012***** Plugin catchall (100. confidence) suggests *****#012#012If you believe that httpd should be allowed getattr access on the lab.html file by default.#012Then you should report this as a bug.#012You can generate a local policy module to allow this access.#012Do#012allow this access for now by executing:#012# ausearch -c 'httpd' --raw | audit2allow -M my-httpd#012# semodule -X 300 -i my-httpd.pp#012
```

4. Run the suggested sealert command. Note the source context, the target objects, the policy, and the enforcing mode.

```
5. [root@serverb ~]# sealert -l 8824e73d-3ab0-4caf-8258-86e8792fee2d
```

6. SELinux is preventing /usr/sbin/httpd from getattr access on the file /lab-content/lab.html.

7.

8. ***** Plugin catchall (100. confidence) suggests *****

9.

- 10.If you believe that httpd should be allowed getattr access on the lab.html file by default.

11.Then you should report this as a bug.

12.You can generate a local policy module to allow this access.

13.Do

14.allow this access for now by executing:

15.# ausearch -c 'httpd' --raw | audit2allow -M my-httpd

16.# semodule -X 300 -i my-httpd.pp

17.

18.

19.Additional Information:

20.Source Context system_u:system_r:httpd_t:s0

21.Target Context unconfined_u:object_r:default_t:s0

```

22.Target Objects          /lab-content/lab.html [ file ]
23.Source                  httpd
24.Source Path             /usr/sbin/httpd
25.Port                   <Unknown>
26.Host                    serverb.lab.example.com
27.Source RPM Packages
28.Target RPM Packages
29.Policy RPM              selinux-policy-3.14.1-59.el8.noarch
30.Selinux Enabled        True
31.Policy Type             targeted
32.Enforcing Mode         Enforcing
33.Host Name               serverb.lab.example.com
34.Platform                Linux serverb.lab.example.com
35.                        4.18.0-67.el8.x86_64 #1 SMP Sat Feb 9
    12:44:00
36.                        UTC 2019 x86_64 x86_64
37.Alert Count             2
38.First Seen              2019-03-28 15:19:46 CET
39.Last Seen               2019-03-28 15:19:46 CET
40.Local ID                8824e73d-3ab0-4caf-8258-86e8792fee2d
41.
42.Raw Audit Messages
43.type=AVC msg=audit(1553782786.213:864): avc: denied { getattr } for
    pid=15606 comm="httpd" path="/lab-content/lab.html" dev="vda1"
    ino=8763212 scontext=system_u:system_r:httpd_t:s0
    tcontext=unconfined_u:object_r:default_t:s0 tclass=file permissive=0
44.
45.

```

```
Hash: httpd,httpd_t,default_t,file,getattr
```

46. The Raw Audit Messages section of the `sealert` command contains information from the `/var/log/audit/audit.log`. Use the `ausearch` command to search the `/var/log/audit/audit.log` file. The `-m` option searches on the message type. The `ts` option searches based on time. This entry identifies the relevant process and file causing the alert. The process is the `httpd` Apache web server, the file is `/lab-content/lab.html`, and the context is `system_r:httpd_t`.

```

47. [root@serverb ~]# ausearch -m AVC -ts recent
48. time->Thu Mar 28 15:19:46 2019

```

```
49. type=PROCTITLE msg=audit(1553782786.213:864):
   proctitle=2F7573722F7362696E2F6874747064002D44464F524547524F554E44

50. type=SYSCALL msg=audit(1553782786.213:864): arch=c000003e syscall=6
   success=no exit=-13 a0=7fb900004930 a1=7fb92dfca8e0 a2=7fb92dfca8e0
   a3=1 items=0 ppid=15491 pid=15606 auid=4294967295 uid=48 gid=48
   euid=48 suid=48 fsuid=48 egid=48 sgid=48 fsgid=48 tty=(none)
   ses=4294967295 comm="httpd" exe="/usr/sbin/httpd"
   subj=system_u:system_r:httpd_t:s0 key=(null)
```

```
type=AVC msg=audit(1553782786.213:864): avc: denied { getattr } for
pid=15606 comm="httpd" path="/lab-content/lab.html" dev="vda1"
ino=8763212 scontext=system_u:system_r:httpd_t:s0
tcontext=unconfined_u:object_r:default_t:s0 tclass=file permissive=0
```

1. Display the SELinux context of the new HTTP document root and the original HTTP document root. Resolve the SELinux issue preventing Apache from serving web content.

1. Use the `ls -dZ` to compare the document root of `/lab-content` and `/var/www/html`.

```
2. [root@serverb ~]# ls -dZ /lab-content /var/www/html
```

```
unconfined_u:object_r:default_t:s0 /lab-content/
system_u:object_r:httpd_sys_content_t:s0 /var/www/html/
```

3. Create a file context rule that sets the default type to `httpd_sys_content_t` for `/lab-content` and all the files below it.

```
4. [root@serverb ~]# semanage fcontext -a \
```

```
-t httpd_sys_content_t '/lab-content(/.*)?'
```

5. Use the `restorecon` command to set the SELinux context for the files in `/lab-content`.

```
[root@serverb ~]# restorecon -R /lab-content/
```

Verify that the SELinux issue has been resolved and Apache is able to serve web content.

Use your web browser to refresh the `http://serverb/lab.html` link. Now you should see some web content.

This is the html file for the SELinux final lab on SERVERB.

1. Exit from `serverb`.

```
2. [root@serverb ~]# exit
```

```
3. logout
```

4. [student@serverb ~]\$ exit
5. logout
6. Connection to serverb closed.

```
[student@workstation ~]$
```

Lab5: Managing Basic Storage

1. New disks are available on serverb. On the first new disk, create a 2 GB GPT partition named backup. Because it may be difficult to set the exact size, a size between 1.8 GB and 2.2 GB is acceptable. Set the correct file-system type on that partition to host an XFS file system.

The password for the student user account on serverb is student. This user has full root access through **sudo**.

1. Use the **ssh** command to log in to serverb as the student user. The systems are configured to use SSH keys for authentication, therefore a password is not required.

2. [student@workstation ~]\$ ssh student@serverb
3. ...output omitted...

```
[student@serverb ~]$
```

4. Because creating partitions and file systems requires root access, use the **sudo -i** command to switch to the root user. If prompted, use student as the password.

5. [student@serverb ~]\$ sudo -i
6. [sudo] password for student: student

```
[root@serverb ~]#
```

7. Use the **lsblk** command to identify the new disks. Those disks should not have any partitions yet.

8. [root@serverb ~]# lsblk
9. NAME MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
10. sr0 11:0 1 1024M 0 rom
11. vda 252:0 0 10G 0 disk
12. └─vda1 252:1 0 10G 0 part /
13. vdb 252:16 0 5G 0 disk
14. vdc 252:32 0 5G 0 disk

```
vdd      252:48    0    5G    0 disk
```

Notice that the first new disk, vdb, does not have any partitions.

15. Confirm that the disk has no label.

```
16. [root@serverb ~]# parted /dev/vdb print
17. Error: /dev/vdb: unrecognised disk label
18. Model: Virtio Block Device (virtblk)
19. Disk /dev/vdb: 5369MB
20. Sector size (logical/physical): 512B/512B
21. Partition Table: unknown
```

Disk Flags:

22. Use **parted** and the **mklabel** subcommand to define the GPT partitioning scheme.

```
23. [root@serverb ~]# parted /dev/vdb mklabel gpt
```

Information: You may need to update /etc/fstab.

24. Create the 2 GB partition. Name it backup and set its type to xfs. Start the partition at sector 2048.

```
25. [root@serverb ~]# parted /dev/vdb mkpart backup xfs 2048s 2GB
```

Information: You may need to update /etc/fstab.

26. Confirm the correct creation of the new partition.

```
27. [root@serverb ~]# parted /dev/vdb print
28. Model: Virtio Block Device (virtblk)
29. Disk /dev/vdb: 5369MB
30. Sector size (logical/physical): 512B/512B
31. Partition Table: gpt
32. Disk Flags:
33.
34. Number  Start      End        Size       File system  Name      Flags
```

1	1049kB	2000MB	1999MB		backup	
---	--------	--------	--------	--	--------	--

35. Run the **udevadm settle** command. This command waits for the system to detect the new partition and to create the `/dev/vdb1` device file. It only returns when it is done.

```
36. [root@serverb ~]# udevadm settle
```

```
[root@serverb ~]#
```

1. Format the 2 GB partition with an XFS file system and persistently mount it at `/backup`.

1. Use the **mkfs.xfs** command to format the `/dev/vdb1` partition.

```
2. [root@serverb ~]# mkfs.xfs /dev/vdb1
```

```
3. meta-data=/dev/vdb1          isize=512    agcount=4,
   agsize=121984 blks
4.          =                   sectsz=512   attr=2, projid32bit=1
5.          =                   crc=1        finobt=1, sparse=1,
   rmapbt=0
6.          =                   reflink=1
7. data      =                   bsize=4096   blocks=487936,
   imaxpct=25
8.          =                   sunit=0      swidth=0 blks
9. naming    =version 2          bsize=4096   ascii-ci=0, ftype=1
10. log      =internal log       bsize=4096   blocks=2560, version=2
11.          =                   sectsz=512   sunit=0 blks,
   lazy-count=1
```

```
realtime =none                extsz=4096   blocks=0, rtextents=0
```

12. Create the `/backup` mount point.

```
13. [root@serverb ~]# mkdir /backup
```

```
[root@serverb ~]#
```

14. Before adding the new file system to `/etc/fstab`, retrieve its UUID.

```
15. [root@serverb ~]# lsblk --fs /dev/vdb1
```

```
16. NAME FSTYPE LABEL UUID                                MOUNTPPOINT
```

```
vdb1 xfs          a3665c6b-4bfb-49b6-a528-74e268b058dd
```

The UUID on your system is probably different.

17. Edit `/etc/fstab` and define the new file system.

```
18. [root@serverb ~]# vim /etc/fstab
```

19. ...output omitted...

```
UUID=a3665c6b-4bfb-49b6-a528-74e268b058dd /backup xfs defaults
0 0
```

20. Force systemd to reread the /etc/fstab file.

```
21. [root@serverb ~]# systemctl daemon-reload
```

```
[root@serverb ~]#
```

22. Manually mount /backup to verify your work. Confirm that the mount is successful.

```
23. [root@serverb ~]# mount /backup
```

```
24. [root@serverb ~]# mount | grep /backup
```

```
/dev/vdb1 on /backup type xfs
(rw,relatime,seclabel,attr2,inode64,noquota)
```

1. On the same new disk, create two 512 MB GPT partitions named swap1 and swap2. A size between 460 MB and 564 MB is acceptable. Set the correct file-system type on those partitions to host swap spaces.
 1. Retrieve the end position of the first partition by displaying the current partition table on /dev/vdb. In the next step, you use that value as the start of the swap1 partition.

```
2. [root@serverb ~]# parted /dev/vdb print
```

```
3. Model: Virtio Block Device (virtblk)
```

```
4. Disk /dev/vdb: 5369MB
```

```
5. Sector size (logical/physical): 512B/512B
```

```
6. Partition Table: gpt
```

```
7. Disk Flags:
```

```
8.
```

```
9. Number  Start      End        Size      File system  Name      Flags
```

```
1      1049kB  2000MB  1999MB  xfs          backup
```

10. Create the first 512 MB partition named swap1. Set its type to linux-swap. Use the end position of the first partition as the starting point. The end position is 2000 MB + 512 MB = 2512 MB

```
11. [root@serverb ~]# parted /dev/vdb mkpart swap1 linux-swap 2000MB
2512M
```

Information: You may need to update /etc/fstab.

12. Create the second 512 MB partition named `swap2`. Set its type to `linux-swap`. Use the end position of the previous partition as the starting point: 2512M. The end position is 2512 MB + 512 MB = 3024 MB

```
13. [root@serverb ~]# parted /dev/vdb mkpart swap2 linux-swap 2512M 3024M
```

Information: You may need to update `/etc/fstab`.

14. Display the partition table to verify your work.

```
15. [root@serverb ~]# parted /dev/vdb print
```

```
16. Model: Virtio Block Device (virtblk)
```

```
17. Disk /dev/vdb: 5369MB
```

```
18. Sector size (logical/physical): 512B/512B
```

```
19. Partition Table: gpt
```

```
20. Disk Flags:
```

```
21.
```

22. Number	Start	End	Size	File system	Name	Flags
23. 1	1049kB	2000MB	1999MB	xfs	backup	
24. 2	2000MB	2512MB	513MB		swap1	swap

3	2512MB	3024MB	512MB		swap2	swap
---	--------	--------	-------	--	-------	------

25. Run the **udevadm settle** command. This command waits for the system to register the new partitions and to create the device files.

```
26. [root@serverb ~]# udevadm settle
```

```
[root@serverb ~]#
```

1. Initialize the two 512 MiB partitions as swap spaces and configure them to activate at boot. Set the swap space on the `swap2` partition to be preferred over the other.

1. Use the **mkswap** command to initialize the swap partitions.

```
2. [root@serverb ~]# mkswap /dev/vdb2
```

```
3. Setting up swapspace version 1, size = 489 MiB (512749568 bytes)
```

```
4. no label, UUID=87976166-4697-47b7-86d1-73a02f0fc803
```

```
5. [root@serverb ~]# mkswap /dev/vdb3
```

```
6. Setting up swapspace version 1, size = 488 MiB (511700992 bytes)
```

```
no label, UUID=4d9b847b-98e0-4d4e-9ef7-dfaaf736b942
```

Take note of the UUIDs of the two swap spaces. You use that information in the next step. If you cannot see the **mkswap** output anymore, use the **lsblk --fs** command to retrieve the UUIDs.

7. Edit `/etc/fstab` and define the new swap spaces. To set the swap space on the `swap2` partition to be preferred over `swap1`, give it a higher priority with the `pri` option.

```
8. [root@serverb ~]# vim /etc/fstab
9. ...output omitted...
10. UUID=a3665c6b-4bfb-49b6-a528-74e268b058dd    /backup xfs    defaults    0
    0
11. UUID=87976166-4697-47b7-86d1-73a02f0fc803    swap    swap    pri=10    0
    0
```

```
UUID=4d9b847b-98e0-4d4e-9ef7-dfaaf736b942    swap    swap    pri=20    0
0
```

12. Force `systemd` to reread the `/etc/fstab` file.

```
13. [root@serverb ~]# systemctl daemon-reload
```

```
[root@serverb ~]#
```

14. Use the **swapon -a** command to activate the new swap spaces. Use the **swapon --show** command to confirm the correct activation of the swap spaces.

```
15. [root@serverb ~]# swapon -a
16. [root@serverb ~]# swapon --show
17. NAME          TYPE          SIZE USED PRI0
18. /dev/vdb2 partition 489M    0B    10
```

```
/dev/vdb3 partition 488M    0B    20
```

1. To verify your work, reboot `serverb`. Confirm that the system automatically mounts the first partition at `/backup`. Also, confirm that the system activates the two swap spaces.

When done, log off from `serverb`.

1. Reboot `serverb`.

```
2. [root@serverb ~]# systemctl reboot
3. [root@serverb ~]#
4. Connection to serverb closed by remote host.
```

5. Connection to serverb closed.

```
[student@workstation ~]$
```

6. Wait a few minutes for serverb to reboot and then log in as the student user.

7. [student@workstation ~]\$ ssh student@serverb

8. ...output omitted...

```
[student@serverb ~]$
```

9. Verify that the system automatically mounts /dev/vdb1 at /backup.

10. [student@serverb ~]\$ mount | grep /backup

```
/dev/vdb1 on /backup type xfs
(rw,relatime,seclabel,attr2,inode64,noquota)
```

11. Use the **swapon --show** command to confirm that the system activates both swap spaces.

12. [student@serverb ~]\$ swapon --show

13. NAME TYPE SIZE USED PRIO

14. /dev/vdb2 partition 489M 0B 10

```
/dev/vdb3 partition 488M 0B 20
```

15. Log off from serverb.

16. [student@serverb ~]\$ exit

17. logout

18. Connection to serverb closed.

```
[student@workstation ~]$
```

Lab7: Managing Logical Volumes

1. Create a 512 MiB partition on /dev/vdb, initialize it as a physical volume, and extend the serverb_01_vg volume group with it.

1. Log in to serverb as the student user.

2. [student@workstation ~]\$ ssh student@serverb

3. ...output omitted...

```
[student@serverb ~]$
```

4. Use the **sudo -i** command to switch to the root user. The password for the student user is student.

```
5. [student@serverb ~]$ sudo -i
```

```
6. [sudo] password for student: student
```

```
[root@serverb ~]#
```

7. Use **parted** to create the 512 MiB partition and set it to type Linux LVM.

```
8. [root@serverb ~]# parted -s /dev/vdb mkpart primary 514MiB 1026MiB
```

```
[root@serverb ~]# parted -s /dev/vdb set 2 lvm on
```

9. Use **udevadm settle** for the system to register the new partition.

```
[root@servera ~]# udevadm settle
```

10. Use **pvccreate** to initialize the partition as a PV.

```
11. [root@serverb ~]# pvcreate /dev/vdb2
```

```
Physical volume "/dev/vdb2" successfully created.
```

12. Use **vgextend** to extend the VG named serverb_01_vg, using the new /dev/vdb2 PV.

```
13. [root@serverb ~]# vgextend serverb_01_vg /dev/vdb2
```

```
Volume group "serverb_01_vg" successfully extended
```

1. Extend the serverb_01_lv logical volume to 768 MiB, including the file system.

1. Use **lvextend** to extend the serverb_01_lv LV to 768 MiB.

```
2. [root@serverb ~]# lvextend -L 768M /dev/serverb_01_vg/serverb_01_lv
```

```
3. Size of logical volume serverb_01_vg/serverb_01_lv changed from 256.00 MiB (64 extents) to 768.00 MiB (192 extents).
```

```
Logical volume serverb_01_vg/serverb_01_lv successfully resized.
```

Note

Alternatively, you could have used the **-L +512M** option to resize the LV.

4. Use **xfs_growfs** to extend the XFS file system to the remainder of the free space on the LV.

```
[root@serverb ~]# xfs_growfs /storage/data1  
6. meta-data=/dev/mapper/serverb_01_vg-serverb_01_lv isize=512  
   agcount=4, agsize=16384 blks
```

...output omitted...

Note

This example shows the **xfs_growfs** step to extend the file system. An alternative would have been to add the **-r** option to the **lvextend** command.

1. In the existing volume group, create a new logical volume called `serverb_02_lv` with a size of 128 MiB. Add an XFS file system and mount it persistently on `/storage/data2`.

1. Use **lvcreate** to create a 128 MiB LV named `serverb_02_lv` from the `serverb_01_vg` VG.

```
[root@serverb ~]# lvcreate -n serverb_02_lv -L 128M serverb_01_vg
```

Logical volume "serverb_02_lv" created

3. Use **mkfs** to place an xfs file system on the `serverb_02_lv` LV. Use the LV device name.

```
[root@serverb ~]# mkfs -t xfs /dev/serverb_01_vg/serverb_02_lv  
5. meta-data=/dev/serverb_01_vg/serverb_02_lv isize=512    agcount=4,  
   agsize=8192 blks
```

...output omitted...

6. Use **mkdir** to create a mount point at `/storage/data2`.

```
[root@serverb ~]# mkdir /storage/data2
```

7. Add the following line to the end of `/etc/fstab` on `serverb`:

```
/dev/serverb_01_vg/serverb_02_lv /storage/data2 xfs defaults 1 2
```

8. Use **systemctl daemon-reload** to update `systemd` with the new `/etc/fstab` configuration.

```
[root@servera ~]# systemctl daemon-reload
```

9. Use **mount** to verify the `/etc/fstab` entry and mount the new `serverb_02_lv` LV device.

```
[root@serverb ~]# mount /storage/data2
```

1. When you are done, reboot your `serverb` machine, then run the command **lab lvm-review grade** from your workstation machine to verify your work.

Wait until `serverb` is completely up and then proceed to run the evaluation.

```
[root@serverb ~]# systemctl reboot
```

Lab8: Implementing Advanced Storage Features

1. From workstation, open an SSH session to `serverb` as student.

2. `[student@workstation ~]$ ssh student@serverb`
3. *...output omitted...*

```
[student@serverb ~]$
```

1. Switch to the root user.

2. `[student@serverb ~]$ sudo -i`
3. `[sudo] password for student: student`

```
[root@serverb ~]#
```

Hide Solution

4. Install the `stratisd` and `stratis-cli` packages using `yum`.

5. `[root@serverb ~]# yum install stratisd stratis-cli`
6. *...output omitted...*
7. `Is this ok [y/N]: y`
8. *...output omitted...*

```
Complete!
```

Hide Solution

9. Start and enable the `stratisd` service using the `systemctl` command.

```
[root@serverb ~]# systemctl enable --now stratisd
```


Hide Solution

10. Create the Stratis pool labpool containing the block device /dev/vdb.

1. Create the Stratis pool labpool using the stratis pool create command.

```
[root@serverb ~]# stratis pool create labpool /dev/vdb
```

2. Verify the availability of labpool using the stratis pool list command.

```
3. [root@serverb ~]# stratis pool list
```

```
4. Name                               Total Physical
```

```
labpool  5 GiB / 37.63 MiB / 4.96 GiB
```

Note the size of the pool in the preceding output.

11. Hide Solution

12. Expand the capacity of labpool using the disk /dev/vdc available in the system.

1. Add the block device /dev/vdc to labpool using the stratis pool add-data command.

```
[root@serverb ~]# stratis pool add-data labpool /dev/vdc
```

2. Verify the size of labpool using the stratis pool list command.

```
3. [root@serverb ~]# stratis pool list
```

```
4. Name                               Total Physical
```

```
labpool  10 GiB / 41.63 MiB / 9.96 GiB
```

The preceding output shows that the size of labpool has increased after a new disk was added to the pool.

5. Use the stratis blockdev list command to list the block devices that are now members of labpool.

```
6. [root@serverb ~]# stratis blockdev list labpool
```

```
7. Pool Name  Device Node  Physical Size  Tier
```

```
8. labpool    /dev/vdb          5 GiB  Data
```

```
labpool      /dev/vdc          5 GiB  Data
```

13. Hide Solution

14. Create a thinly provisioned file system named `labfs` in the `labpool` pool. Mount this file system on `/labstratisvol` so that it persists across reboots. Create a file named `labfile1` that contains the text `Hello World!` on the `labfs` file system. Don't forget to use

the `x-systemd.requires=stratisd.service` mount option in `/etc/fstab`.

1. Create the thinly provisioned file system `labfs` in `labpool` using the `stratis filesystem create` command. It may take up to a minute for the command to complete.

```
[root@serverb ~]# stratis filesystem create labpool labfs
```

2. Verify the availability of `labfs` using the `stratis filesystem list` command.

```
3. [root@serverb ~]# stratis filesystem list
```

4. Pool Name	Name	Used	Created	Device
--------------	------	------	---------	--------

labpool	labfs	546 MiB	Mar 29 2019 07:48	/stratis/labpool/labfs 9825...d6ca
---------	-------	---------	-------------------	---------------------------------------

Note the current usage of `labfs`. This usage of the file system increases on-demand in the following steps.

5. Determine the UUID of `labfs` using the `lsblk` command.

```
6. [root@serverb ~]# lsblk --output=UUID /stratis/labpool/labfs
```

```
7. UUID
```

```
9825e289-fb08-4852-8290-44d1b8f0d6ca
```

8. Edit `/etc/fstab` so that the thinly provisioned file system `labfs` is mounted at boot time. Use the UUID you determined in the preceding step. The following shows the line you should add to `/etc/fstab`. You can use the `vi /etc/fstab` command to edit the file.

```
UUID=9825...d6ca /labstratisvol xfs  
defaults,x-systemd.requires=stratisd.service 0 0
```

9. Create a directory named `/labstratisvol` using the `mkdir` command.

```
[root@serverb ~]# mkdir /labstratisvol
```

10. Mount the thinly provisioned file system `labfs` using the `mount` command to confirm that the `/etc/fstab` file contains the appropriate entries.

```
[root@serverb ~]# mount /labstratisvol
```

If the preceding command produces any errors, revisit the `/etc/fstab` file and ensure that it contains the appropriate entries.

11. Create a text file named `/labstratisvol/labfile1` using the `echo` command.

```
[root@serverb ~]# echo "Hello World!" > /labstratisvol/labfile1
```

15. Hide Solution

16. Verify that the thinly provisioned file system `labfs` dynamically grows as the data on the file system grows by adding a 2 GiB `labfile2` to the filesystem.

1. View the current usage of `labfs` using the `stratis filesystem list` command.

```
2. [root@serverb ~]# stratis filesystem list
```

3. Pool Name	Name	Used	Created	Device
UUID				

labpool	labfs	546 MiB	Mar 29 2019 07:48	/stratis/labpool/labfs 9825...d6ca
---------	-------	---------	-------------------	---------------------------------------

4. Create a 2 GiB file in `labfs` using the `dd` command. It may take up to a minute for the command to complete.

```
[root@serverb ~]# dd if=/dev/urandom of=/labstratisvol/labfile2 bs=1M  
count=2048
```

5. Verify that the usage of `labfs` has increased, using the `stratis filesystem list` command.

```
6. [root@serverb ~]# stratis filesystem list
```

7. Pool Name	Name	Used	Created	Device
Device			UUID	
8. labpool	labfs	2.53 GiB	Mar 29 2019 07:48	/stratis/labpool/labfs 9825...d6ca

17. Hide Solution

18. Create a snapshot named `labfs-snap` of the `labfs` file system. The snapshot allows you to access any file that is deleted from `labfs`.

1. Create a snapshot of `labfs` using the `stratis filesystem snapshot` command. It may take up to a minute for the command to complete.

```
2. [root@serverb ~]# stratis filesystem snapshot labpool \
```

```
labfs labfs-snap
```

3. Verify the availability of the snapshot using the `stratis filesystem list` command.

```
4. [root@serverb ~]# stratis filesystem list
```

```
5. ...output omitted...
```

```
labpool labfs-snap 2.53 GiB Mar 29 2019 10:28
/stratis/labpool/labfs-snap 291d...8a16
```

6. Remove the `/labstratisvol/labfile1` file.

```
7. [root@serverb ~]# rm /labstratisvol/labfile1
```

```
rm: remove regular file '/labstratisvol/labfile1'? y
```

8. Create the `/labstratisvol-snap` directory using the `mkdir` command.

```
[root@serverb ~]# mkdir /labstratisvol-snap
```

9. Mount the snapshot `labfs-snap` on `/labstratisvol-snap` using the `mount` command.

```
10. [root@serverb ~]# mount /stratis/labpool/labfs-snap \
```

```
/labstratisvol-snap
```

11. Confirm that you can still access the file you deleted from `labfs` using the snapshot `labfs-snap`.

```
12. [root@serverb ~]# cat /labstratisvol-snap/labfile1
```

```
Hello World!
```

19. Hide Solution

20. Create the VDO volume `labvdo`, with the device `/dev/vdd`. Set its logical size to 50 GB.

1. Create the volume `labvdo` using the `vdo create` command.

```
2. [root@serverb ~]# vdo create --name=labvdo --device=/dev/vdd
--vdoLogicalSize=50G
```

```
...output omitted...
```

3. Verify the availability of the volume `labvdo` using the `vdo list` command.

```
[root@serverb ~]# vdo list
```

```
labvdo
```

21. Hide Solution

22. Mount the volume `labvdo` on `/labvdovol` with the `xfs` file system so that it persists across reboots. Don't forget to use the `x-systemd.requires=vdo.service` mount option in `/etc/fstab`.

1. Format the `labvdo` volume with the `xfs` file system using the `mkfs` command.

```
[root@serverb ~]# mkfs.xfs -K /dev/mapper/labvdo
```

```
...output omitted...
```

3. Use the `udevadm` command to register the new device node.

```
[root@serverb ~]# udevadm settle
```

4. Create the `/labvdovol` directory using the `mkdir` command.

```
[root@serverb ~]# mkdir /labvdovol
```

5. Determine the UUID of `labvdo` using the `lsblk` command.

```
[root@serverb ~]# lsblk --output=UUID /dev/mapper/labvdo
```

7. UUID

```
ef8cce71-228a-478d-883d-5732176b39b1
```

8. Edit `/etc/fstab` so that `labvdo` is mounted at boot time. Use the UUID of the volume you determined in the preceding step. The following shows the line you should add to `/etc/fstab`. You can use the `vi /etc/fstab` command to edit the file.

```
UUID=ef8c...39b1 /labvdovol xfs
defaults,x-systemd.requires=vdo.service 0 0
```

9. Mount the `labvdo` volume using the `mount` command to confirm that the `/etc/fstab` file contains the appropriate entries.

```
[root@serverb ~]# mount /labvdovol
```

If the preceding command produces any errors, revisit the `/etc/fstab` file and ensure that it contains the appropriate entries.

23. Hide Solution

24. Create three copies of the file named `/root/install.img` on the volume `labvdo`. Compare the statistics of the volume to verify the data deduplication and compression happening on the volume.

1. View the initial statistics and status of the volume using the `vdostats` command.

```
2. [root@serverb ~]# vdostats --human-readable
```

3. Device	Size	Used	Available	Use%	Space saving%
-----------	------	------	-----------	------	---------------

<code>/dev/mapper/labvdo</code>	5.0G	3.0G	2.0G	60%	99%
---------------------------------	------	------	------	-----	-----

Notice that 3 GB of the volume is already used because when created, the VDO volume reserves 3-4 GB for itself. Also note that the value 99% in the `Space saving%` field indicates that you have not created any content so far in the volume, contributing to all of the saved volume space.

4. Copy `/root/install.img` to `/labvdovol/install.img.1` and verify the statistics of the volume. It may take up to a minute to copy the file.

```
5. [root@serverb ~]# cp /root/install.img /labvdovol/install.img.1
```

```
6. [root@serverb ~]# vdostats --human-readable
```

7. Device	Size	Used	Available	Use%	Space saving%
-----------	------	------	-----------	------	---------------

<code>/dev/mapper/labvdo</code>	5.0G	3.4G	1.6G	68%	5%
---------------------------------	------	------	------	-----	----

Notice that the value of the `Used` field increased from 3.0G to 3.4G because you copied a file in the volume, and that occupies some space. Also, notice that the value of `Space saving%` field decreased from 99% to 5% because initially there was no content in the volume, contributing to the low volume space utilization and high volume space saving until you created a file in there. The volume space saving is quite low because you created a unique copy of the file in the volume and there is nothing to deduplicate.

8. Copy `/root/install.img` to `/labvdovol/install.img.2` and verify the statistics of the volume. It may take up to a minute to copy the file.

```
9. [root@serverb ~]# cp /root/install.img /labvdovol/install.img.2
```

```
10. [root@serverb ~]# vdostats --human-readable
```

11. Device	Size	Used	Available	Use%	Space saving%
/dev/mapper/labvdo 51%	5.0G	3.4G	1.6G	68%	

Notice that the used volume space did not change. Instead, the percentage of the saved volume space increased, proving that the data deduplication occurred to reduce the space consumption for the redundant copies of the same file. The value of Space saving% in the preceding output may vary on your system.

25. Hide Solution

26. Reboot serverb. Verify that your labvdo volume is mounted on /labvdovol after the system starts back up.

1. Reboot the serverb machine.

```
[root@serverb ~]# systemctl reboot
```

Note

Note: If on a reboot, serverb does not boot to a regular login prompt but instead has "Give root password for maintenance (or press Control-D to continue):" you likely made a mistake in /etc/fstab. After providing the root password of redhat, you will need to remount the root file system as read-write with:

```
[root@serverb ~]# mount -o remount,rw /
```

Verify that /etc/fstab is configured correctly as specified in the solutions. Pay special attention to the mount options for the lines related to /labstratisvol and /labvdovol.

Lab9: Accessing Network-Attached Storage

1. Log in to servera and install the required packages.

1. Log in to servera as the student user.

```
2. [student@workstation ~]$ ssh student@servera
3. ...output omitted...
```

```
[student@servera ~]$
```

4. Use the **sudo -i** command to switch to the root user. The password for the student user is student.

```
5. [student@servera ~]$ sudo -i
```

```
6. [sudo] password for student: student
```

```
[root@servera ~]#
```

7. Install the autofs package.

```
8. [root@servera ~]# yum install autofs
```

```
9. ...output omitted...
```

```
10. Is this ok [y/N]: y
```

```
...output omitted...
```

2. Hide Solution

3. Configure an automounter indirect map on servera using shares from serverb. Create an indirect map using files named /etc/auto.master.d/shares.autofs for the master map and /etc/auto.shares for the mapping file. Use the /remote directory as the main mount point on servera. Reboot servera to determine if the autofs service starts automatically.

1. Test the NFS server before proceeding to configure the automounter.

```
2. [root@servera ~]# mount -t nfs serverb.lab.example.com:/shares /mnt
```

```
3. [root@servera ~]# ls -l /mnt
```

```
4. total 0
```

```
5. drwxrwx---. 2 root managers  25 Apr  4 01:13 management
```

```
6. drwxrwx---. 2 root operators  25 Apr  4 01:13 operation
```

```
7. drwxrwx---. 2 root production 25 Apr  4 01:13 production
```

```
[root@servera ~]# umount /mnt
```

8. Create a master map file named /etc/auto.master.d/shares.autofs, insert the following content, and save the changes.

```
9. [root@servera ~]# vim /etc/auto.master.d/shares.autofs
```

```
/remote    /etc/auto.shares
```

10. Create an indirect map file named /etc/auto.shares, insert the following content, and save the changes.

```
11. [root@servera ~]# vim /etc/auto.shares
```

```
*    -rw,sync,fstype=nfs4    serverb.lab.example.com:/shares/&
```

12. Start and enable the autofs service on servera.


```
13. [root@servera ~]# systemctl enable --now autofs
```

```
Created symlink  
/etc/systemd/system/multi-user.target.wants/autofs.service →  
/usr/lib/systemd/system/autofs.service.
```

14. Reboot the servera machine.

```
[root@servera ~]# systemctl reboot
```

4. Hide Solution

5. Test the autofs configuration with the various users. When done, log off from servera.

1. After the servera machine has finished booting, log in to servera as the student user.

```
2. [student@workstation ~]$ ssh student@servera
```

```
3. ...output omitted...
```

```
[student@servera ~]$
```

4. Use the **su - manager1** command to switch to the manager1 user and test access.

```
5. [student@servera ~]$ su - manager1
```

```
6. Password: redhat
```

```
7. [manager1@servera ~]$ ls -l /remote/management/
```

```
8. total 4
```

```
9. -rw-r--r--. 1 root managers 46 Apr  4 01:13 Welcome.txt
```

```
10. [manager1@servera ~]$ cat /remote/management/Welcome.txt
```

```
11. ###Welcome to Management Folder on SERVERB###
```

```
12. [manager1@servera ~]$ echo TEST1 > /remote/management/Test.txt
```

```
13. [manager1@servera ~]$ cat /remote/management/Test.txt
```

```
14. TEST1
```

```
15. [manager1@servera ~]$ ls -l /remote/operation/
```

```
16. ls: cannot open directory '/remote/operation/': Permission denied
```

```
17. [manager1@servera ~]$ ls -l /remote/production/
```

```
18. ls: cannot open directory '/remote/production/': Permission denied
```

```
19. [manager1@servera ~]$ exit
```

```
20. logout
```

```
21. [student@servera ~]$
```

22. Switch to the dbuser1 user and test access.

```
23. [student@servera ~]$ su - dbuser1
24. Password: redhat
25. [dbuser1@servera ~]$ ls -l /remote/production/
26. total 4
27. -rw-r--r--. 1 root production 46 Apr  4 01:13 Welcome.txt
28. [dbuser1@servera ~]$ cat /remote/production/Welcome.txt
29. ###Welcome to Production Folder on SERVERB###
30. [dbuser1@servera ~]$ echo TEST2 > /remote/production/Test.txt
31. [dbuser1@servera ~]$ cat /remote/production/Test.txt
32. TEST2
33. [dbuser1@servera ~]$ ls -l /remote/operation/
34. ls: cannot open directory '/remote/operation/': Permission denied
35. [dbuser1@servera ~]$ ls -l /remote/management/
36. ls: cannot open directory '/remote/management/': Permission denied
37. [dbuser1@servera ~]$ exit
38. logout
39. [student@servera ~]$
```

40. Switch to the contractor1 user and test access.

```
41. [student@servera ~]$ su - contractor1
42. Password: redhat
43. [contractor1@servera ~]$ ls -l /remote/operation/
44. total 4
45. -rw-r--r--. 1 root operators 45 Apr  4 01:13 Welcome.txt
46. [contractor1@servera ~]$ cat /remote/operation/Welcome.txt
47. ###Welcome to Operation Folder on SERVERB###
48. [contractor1@servera ~]$ echo TEST3 > /remote/operation/Test.txt
49. [contractor1@servera ~]$ cat /remote/operation/Test.txt
50. TEST3
51. [contractor1@servera ~]$ ls -l /remote/management/
52. ls: cannot open directory '/remote/management/': Permission denied
53. [contractor1@servera ~]$ ls -l /remote/production/
54. ls: cannot open directory '/remote/production/': Permission denied
55. [contractor1@servera ~]$ exit
56. logout
```

```
57. [student@servera ~]$
```

58. Explore the **mount** options for the NFS automounted share.

```
59. [student@servera ~]$ mount | grep nfs
```

```
60. rpc_pipefs on /var/lib/nfs/rpc_pipefs type rpc_pipefs (rw,relatime)
```

```
61. serverb.lab.example.com:/shares/management on /remote/management type  
nfs4
```

```
62. (rw,relatime,vers=4.2,rsz=262144,wsz=262144,namlen=255,
```

```
63. sync,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=172.25.250.10,
```

```
64. local_lock=none,addr=172.25.250.11)
```

```
65. serverb.lab.example.com:/shares/operation on /remote/operation type  
nfs4
```

```
66. (rw,relatime,vers=4.2,rsz=262144,wsz=262144,namlen=255,
```

```
67. sync,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=172.25.250.10,
```

```
68. local_lock=none,addr=172.25.250.11)
```

```
69. serverb.lab.example.com:/shares/production on /remote/production type  
nfs4
```

```
70. (rw,relatime,vers=4.2,rsz=262144,wsz=262144,namlen=255,
```

```
71. sync,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=172.25.250.10,
```

```
local_lock=none,addr=172.25.250.11)
```

72. Log off from servera.

```
73. [student@servera ~]$ exit
```

```
74. logout
```

```
[student@workstation ~]$
```

Lab10: Controlling the Boot Process

1. On serverb, reset the root password to redhat.

Locate the icon for the serverb console, as appropriate for your classroom environment. Work from that console.

1. Send a **Ctrl+Alt+Del** to your system using the relevant button or menu entry.
2. When the boot-loader menu appears, press any key to interrupt the countdown, except **Enter**.
3. Use the cursor keys to highlight the default boot loader entry.
4. Press **e** to edit the current entry.

5. Use the cursor keys to navigate to the line that starts with `linux`.
6. Press **End** to move the cursor to the end of the line.
7. Append `rd.break` to the end of the line.
8. Press **Ctrl+x** to boot using the modified configuration.
9. At the `switch_root` prompt, remount the `/sysroot` file system read/write, then use **chroot** to go into a **chroot** jail at `/sysroot`.

```
10. switch_root:/# mount -o remount,rw /sysroot
```

```
switch_root:/# chroot /sysroot
```

11. Set the root password to `redhat`.

```
12. sh-4.4# passwd root
```

```
13. Changing password for user root.
```

```
14. New password: redhat
```

```
15. BAD PASSWORD: The password is shorter than 8 characters
```

```
16. Retype new password: redhat
```

```
passwd: all authentication tokens updated successfully.
```

17. Configure the system to automatically perform a full SELinux relabel after boot.

```
sh-4.4# touch /.autorelabel
```

18. Type **exit** twice to continue booting your system. The system fails to boot because of an issue you resolve in the next step.

Hide Solution

2. The system fails to boot. A start job does not seem to complete. From the console, fix the issue.
 1. Boot the system into emergency mode. To do so, reboot serverb by sending a **Ctrl+Alt+Del** to your system using the relevant button or menu entry.
 2. When the boot-loader menu appears, press any key to interrupt the countdown, except **Enter**.
 3. Use the cursor keys to highlight the default boot loader entry.
 4. Press **e** to edit the current entry.
 5. Use the cursor keys to navigate to the line that starts with `linux`.
 6. Press **End** to move the cursor to the end of the line.
 7. Append `systemd.unit=emergency.target` to the end of the line.
 8. Press **Ctrl+x** to boot using the modified configuration.
 9. Log in to emergency mode. The root password is `redhat`.

```
10. Give root password for maintenance
```

11. (or press Control-D to continue): **redhat**

```
[root@serverb ~]#
```

12. Remount the / file system read/write.

```
[root@serverb ~]# mount -o remount,rw /
```

13. Use the **mount -a** command to attempt to mount all the other file systems.

```
14. [root@serverb ~]# mount -a
```

```
mount: /olddata: can't find
UUID=4d5c85a5-8921-4a06-8aff-80567e9689bc.
```

15. Edit /etc/fstab to remove or comment out the incorrect line.

```
16. [root@serverb ~]# vim /etc/fstab
```

17. ...*output omitted*...

```
#UUID=4d5c85a5-8921-4a06-8aff-80567e9689bc /olddata xfs defaults
0 0
```

18. Update systemd for the system to register the new /etc/fstab configuration.

```
19. [root@serverb ~]# systemctl daemon-reload
```

```
[root@serverb ~]#
```

20. Verify that your /etc/fstab is now correct by attempting to mount all entries.

```
21. [root@serverb ~]# mount -a
```

```
[root@serverb ~]#
```

22. Reboot the system and wait for the boot to complete. Because you created the /.autorelabel file in the first step, after setting the root password, the system runs an SELinux relabel, then reboots again by itself. The system should now boot normally.

```
[root@serverb ~]# systemctl reboot
```

3. Hide Solution

4. Change the default `systemd` target on `serverb` for the system to automatically start a graphical interface when it boots.

No graphical interface is installed yet on `serverb`. For this exercise, only set the default target and do not install the packages.

1. Log in to `serverb` as the `root` user. Use `redhat` as the password.
2. Use the **`systemctl set-default`** command to set `graphical.target` as the default target.

```
[root@serverb ~]# systemctl set-default graphical.target
```

3. Use the **`systemctl get-default`** command to verify your work.

```
4. [root@serverb ~]# systemctl get-default
```

```
graphical.target
```

5. Log off from `serverb`.

```
[root@serverb ~]# exit
```

Lab11: Managing Network Security

1. From workstation, test access to the default web server at `http://serverb.lab.example.com` and to the virtual host at `http://serverb.lab.example.com:1001`.
 1. Test access to the `http://serverb.lab.example.com` web server. The test currently fails. Ultimately, the web server should return `SERVER B`.

```
2. [student@workstation ~]$ curl http://serverb.lab.example.com
3. curl: (7) Failed to connect to serverb.lab.example.com port 80:
   Connection refused
```

4. Test access to the `http://serverb.lab.example.com:1001` virtual host. The test currently fails. Ultimately, the virtual host should return `VHOST 1`.

```
5. [student@workstation ~]$ curl http://serverb.lab.example.com:1001
6. curl: (7) Failed to connect to serverb.lab.example.com port 1001: No
   route to host
```

Hide Solution

2. Log in to `serverb` to determine what is preventing access to the web servers.

1. From workstation, open an SSH session to serverb as student user. The systems are configured to use SSH keys for authentication, so a password is not required.

```
2. [student@workstation ~]$ ssh student@serverb
3. ...output omitted...
4. [student@serverb ~]$
```

5. Determine whether the httpd service is active.

```
6. [student@serverb ~]$ systemctl is-active httpd
7. inactive
```

8. Enable and start the httpd service. The httpd service fails to start.

```
9. [student@serverb ~]$ sudo systemctl enable --now httpd
10. [sudo] password for student: student
11. Created symlink
    /etc/systemd/system/multi-user.target.wants/httpd.service →
    /usr/lib/systemd/system/httpd.service.
12. Job for httpd.service failed because the control process exited with
    error code.
13. See "systemctl status httpd.service" and "journalctl -xe" for
    details.
```

14. Investigate the reasons why the httpd.service service failed to start.

```
15. [student@serverb ~]$ systemctl status httpd.service
16. ● httpd.service - The Apache HTTP Server
17.    Loaded: loaded (/usr/lib/systemd/system/httpd.service; enabled;
    vendor preset: disabled)
18.    Active: failed (Result: exit-code) since Thu 2019-04-11 19:25:36
    CDT; 19s ago
19.       Docs: man:httpd.service(8)
20.   Process: 9615 ExecStart=/usr/sbin/httpd $OPTIONS -DFOREGROUND
    (code=exited, status=1/FAILURE)
21.  Main PID: 9615 (code=exited, status=1/FAILURE)
22.    Status: "Reading configuration..."
23.
24. Apr 11 19:25:36 serverb.lab.example.com systemd[1]: Starting The
    Apache HTTP Server...
25. Apr 11 19:25:36 serverb.lab.example.com httpd[9615]: (13)Permission
    denied: AH00072: make_sock: could not bind to address [::]:1001
26. Apr 11 19:25:36 serverb.lab.example.com httpd[9615]: (13)Permission
    denied: AH00072: make_sock: could not bind to address 0.0.0.0:1001
```

```

27. Apr 11 19:25:36 serverb.lab.example.com httpd[9615]: no listening
    sockets available, shutting down
28. Apr 11 19:25:36 serverb.lab.example.com httpd[9615]: AH00015: Unable
    to open logs
29. Apr 11 19:25:36 serverb.lab.example.com systemd[1]: httpd.service:
    Main process exited, code=exited, status=1/FAILURE
30. Apr 11 19:25:36 serverb.lab.example.com systemd[1]: httpd.service:
    Failed with result 'exit-code'.
31. Apr 11 19:25:36 serverb.lab.example.com systemd[1]: Failed to start
    The Apache HTTP Server.

```

32. Use the **sealert** command to check whether SELinux is blocking the httpd service from binding to port 1001/TCP.

```

33. [student@serverb ~]$ sudo sealert -a /var/log/audit/audit.log
34. 100% done
35. found 1 alerts in /var/log/audit/audit.log
36. -----
    -----
37.
38. SELinux is preventing /usr/sbin/httpd from name_bind access on the
    tcp_socket port 1001.
39.
40. ***** Plugin bind_ports (99.5 confidence) suggests
    *****
41.
42. If you want to allow /usr/sbin/httpd to bind to network port 1001
43. Then you need to modify the port type.
44. Do
45. # semanage port -a -t PORT_TYPE -p tcp 1001
46.     where PORT_TYPE is one of the following: http_cache_port_t,
    http_port_t, jboss_management_port_t, jboss_messaging_port_t,
    ntop_port_t, puppet_port_t.
47.
48. ***** Plugin catchall (1.49 confidence) suggests
    *****
49.
50. ...output omitted...

```

Hide Solution

3. Configure SELinux to allow the httpd service to listen on port 1001/TCP.

1. Use the **semanage** command to find the correct port type.

```

2. [student@serverb ~]$ sudo semanage port -l | grep 'http'

```



```
3. http_cache_port_t      tcp  8080, 8118, 8123, 10001-10010
4. http_cache_port_t      udp  3130
5. http_port_t            tcp  80, 81, 443, 488, 8008, 8009, 8443, 9000
6. pegasus_http_port_t    tcp  5988
7. pegasus_https_port_t   tcp  5989
```

8. Use the **semanage** command to bind port 1001/TCP to the http_port_t type.

```
9. [student@serverb ~]$ sudo semanage port -a -t http_port_t -p tcp 1001
10. [student@serverb ~]$
```

11. Confirm that port 1001/TCP is bound to the http_port_t port type.

```
12. [student@serverb ~]$ sudo semanage port -l | grep '^http_port_t'
13. http_port_t            tcp  1001, 80, 81, 443, 488, 8008, 8009, 8443,
    9000
```

14. Enable and start the httpd service.

```
15. [student@serverb ~]$ sudo systemctl enable --now httpd
```

16. Verify the running state of the httpd service.

```
17. [student@serverb ~]$ systemctl is-active httpd; systemctl is-enabled httpd
18. active
19. enabled
```

20. Exit from serverb.

```
21. [student@serverb ~]$ exit
22. logout
23. Connection to serverb closed.
24. [student@workstation ~]$
```

Hide Solution

4. From workstation, test access to the default web server at `http://serverb.lab.example.com` and to the virtual host at `http://serverb.lab.example.com:1001`.
 1. Test access to the `http://serverb.lab.example.com` web server. The web server should return `SERVER B`.

```
2. [student@workstation ~]$ curl http://serverb.lab.example.com
3. SERVER B
```

4. Test access to the `http://serverb.lab.example.com:1001` virtual host. The test continues to fail.

```
5. [student@workstation ~]$ curl http://serverb.lab.example.com:1001
6. curl: (7) Failed to connect to serverb.lab.example.com port 1001: No
   route to host
```

Hide Solution

5. Log in to serverb to determine whether the correct ports are assigned to the firewall.

1. From workstation, log in to serverb as the student user.

```
2. [student@workstation ~]$ ssh student@serverb
3. ...output omitted...
4. [student@serverb ~]$
```

5. Verify that the default firewall zone is set to public.

```
6. [student@serverb ~]$ firewall-cmd --get-default-zone
```

```
public
```

7. If the previous step did not return `public` as the default zone, correct it with the following command:

```
[student@serverb ~]$ sudo firewall-cmd --set-default-zone public
```

8. Determine the open ports listed in the `public` network zone.

```
9. [student@serverb ~]$ sudo firewall-cmd --permanent --zone=public
   --list-all
10. [sudo] password for student: student
11. public
12.   target: default
13.   icmp-block-inversion: no
14.   interfaces:
15.   sources:
16.   services: cockpit dhcpv6-client http ssh
17.   ports:
18.   protocols:
19.   masquerade: no
20.   forward-ports:
21.   source-ports:
22.   icmp-blocks:
```

23. rich rules:

Hide Solution

6. Add port 1001/TCP to the permanent configuration for the public network zone. Confirm your configuration.

1. Add port 1001/TCP to the public network zone.

```
2. [student@serverb ~]$ sudo firewall-cmd --permanent --zone=public \  
3. --add-port=1001/tcp  
4. success
```

5. Reload the firewall configuration.

```
6. [student@serverb ~]$ sudo firewall-cmd --reload  
7. success
```

8. Confirm your configuration.

```
9. [student@serverb ~]$ sudo firewall-cmd --permanent --zone=public  
--list-all  
10. public  
11. target: default  
12. icmp-block-inversion: no  
13. interfaces:  
14. sources:  
15. services: cockpit dhcpv6-client http ssh  
16. ports: 1001/tcp  
17. protocols:  
18. masquerade: no  
19. forward-ports:  
20. source-ports:  
21. icmp-blocks:  
22. rich rules:
```

23. Exit from serverb.

```
24. [student@serverb ~]$ exit  
25. logout  
26. Connection to serverb closed.  
27. [student@workstation ~]$
```

Hide Solution

7. From workstation, confirm that the default web server at serverb.lab.example.com returns SERVER B and the virtual host at serverb.lab.example.com:1001 returns VHOST 1.
 1. Test access to the http://serverb.lab.example.com web server.

```
2. [student@workstation ~]$ curl http://serverb.lab.example.com
3. SERVER B
```

4. Test access to the http://serverb.lab.example.com:1001 virtual host.

```
5. [student@workstation ~]$ curl http://serverb.lab.example.com:1001
6. VHOST 1
```

Hide Solution

Lab12: Installing Red Hat Enterprise Linux

1. On serverb, copy /root/anaconda-ks.cfg to /home/student/kickstart.cfg, so the student user can edit it.
 1. Use the **ssh** command to log in to serverb as the student user.

```
2. [student@workstation ~]$ ssh student@serverb
3. ...output omitted...
```

```
[student@serverb ~]$
```

4. Copy /root/anaconda-ks.cfg on serverb to a file called /home/student/kickstart.cfg so that student can edit. Use the **sudo cat /root/anaconda-ks.cfg > ~/kickstart.cfg** command to copy the contents of /root/anaconda-ks.cfg to /home/student/kickstart.cfg. If **sudo** prompts for the password of the student user, use student as the password.

```
5. [student@serverb ~]$ sudo cat /root/anaconda-ks.cfg > ~/kickstart.cfg
```

```
[sudo] password for student: student
```

2. Hide Solution
3. Make the following changes to /home/student/kickstart.cfg.
 - Comment out the **reboot** command.
 - Comment out the **repo** command for the BaseOS repository. Modify the **repo** command for the AppStream repository to point to http://classroom.example.com/content/rhel8.2/x86_64/dvd/AppStream/. The repository name should be set to appstream.

- Change the **url** command to use `http://classroom.example.com/content/rhel8.2/x86_64/dvd/` as the installation source.
- Comment out the **network** command.
- Change the **rootpw** command to use `plaintext` and set the root password to `redhat`.
- Delete the line that uses the **auth** command and add the **authselect select sssd** line to set the `sss`d service as the identity and authentication source.
- Simplify the **services** command so that only the `kdump` and `rhsmcertd` services are disabled. Leave only the `sshd`, `rngd`, and `chronyd` enabled.
- Add the **autopart** command. The **part** commands should be commented out.
- Simplify the `%post` section so that it only runs a script to append the text `Kickstarted on DATE` to the end of the `/etc/issue` file. `DATE` is variable information and should be generated by the script using the **date** command with no additional options.
- Simplify the `%package` section as follows: include the `@core`, `chrony`, `dracut-config-generic`, `dracut-norescue`, `firewalld`, `g`, `rub2`, `kernel`, `rsync`, `tar`, and `httpd` packages. Ensure that the `plymouth` package is not installed.

0. Comment out the reboot directive:

```
#reboot
```

1. The **repo** command is found twice in `kickstart.cfg`. Comment out the **repo** command for the BaseOS repository. Modify the **repo** command for the AppStream repository to point to the classroom's AppStream repository:

```
2. #repo --name="koji-override-0"
   --baseurl=http://download-node-02.eng.bos.redhat.com/rhel-8/devel/candidate-trees/RHEL-8/RHEL-8.2.0-updates-20200423.0/compose/BaseOS/x86_64/os
```

```
repo --name="appstream"
--baseurl=http://classroom.example.com/content/rhel8.2/x86_64/dvd/AppStream/
```

3. Change the **url** command to specify the HTTP installation source media used in the classroom:

```
url --url="http://classroom.example.com/content/rhel8.2/x86_64/dvd/"
```

4. Comment out the **network** command:

```
#network --bootproto=dhcp --device=link --activate
```

5. Set the root password to redhat. Change the line that starts with **rootpw** to:

```
rootpw --plaintext redhat
```

6. Delete the line that uses the **auth** command and add the **authselect select sssd** line to set the sssd service as the identity and authentication source.

```
authselect select sssd
```

7. Simplify the **services** command to look exactly like the following:

```
services --disabled="kdump,rhsmcertd" --enabled="sshd,rngd,chronyd"
```

8. Comment out the **part** commands. Add the **autopart** command:

```
9. # Disk partitioning information
10. #part biosboot --fstype="biosboot" --size=1
11. #part /boot/efi --fstype="efi" --size=100 --fsoptions="..."
12. #part / --fstype="xfs" --size=10137 --label=root
```

```
autopart
```

13. Delete all content between the **%post** section and its **%end**. Add the following line: **echo "Kickstarted on \$(date)" >> /etc/issue**

The entire **%post** section should look like this.

```
%post --erroronfail
echo "Kickstarted on $(date)" >> /etc/issue
%end
```

14. Simplify the package specification to look exactly like the following:

```
15. %packages
16. @core
17. chrony
18. dracut-config-generic
19. dracut-norescue
20. firewallld
```

```
21. grub2
22. kernel
23. rsync
24. tar
25. httpd
26. -plymouth
```

```
%end
```

2. Hide Solution
3. Validate the syntax of kickstart.cfg.
 0. Use the **ksvalidator** command to check the Kickstart file for syntax errors.

```
[student@serverb ~]$ ksvalidator kickstart.cfg
```

4. Hide Solution
5. Make the /home/student/kickstart.cfg file available at <http://serverb.lab.example.com/ks-config/kickstart.cfg>
 0. Copy kickstart.cfg to the /var/www/html/ks-config/ directory.

```
[student@serverb ~]$ sudo cp ~/kickstart.cfg /var/www/html/ks-config
```

6. Hide Solution
7. Return to the workstation system to check your work.
 0. Exit from serverb.

```
1. [student@serverb ~]$ exit
2. logout
3. Connection to serverb closed.
```

```
[student@workstation ~]$
```

Lab13: Running Containers

1. On serverb, install the container tools. Log in to serverb as the student user, and then use the sudo command. The password for the student user is student.
 1. Use the ssh command to log in to serverb as the student user. The systems are configured to use SSH keys for authentication, so a password is not required.

```
2. [student@workstation ~]$ ssh student@serverb
```

3. ...output omitted...

```
[student@serverb ~]$
```

4. Install the container-tools Yum module using the yum command.

```
5. [student@serverb ~]$ sudo yum module install container-tools
```

```
6. [sudo] password for student: student
```

7. ...output omitted...

```
8. Is this ok [y/N]: y
```

9. ...output omitted...

Complete!

2. Hide Solution

3. The container image registry at `registry.lab.example.com` stores the `rhel8/mariadb-103` image with several tags. On `serverb`, as the `podsvc` user, list those tags and take note of the tag with the *lowest* version number. You will use that image tag to start a container later in this exercise.

The password for the `podsvc` user is `redhat`. To query the `registry.lab.example.com` registry, use the `admin` account with `redhat321` for the password.

1. Exit from the student account on `serverb`.

```
2. [student@serverb ~]$ exit
```

```
3. logout
```

```
4. Connection to serverb closed.
```

```
[student@workstation ~]$
```

5. Use the `ssh` command to log in to `serverb` as the `podsvc` user. The systems are configured to use SSH keys for authentication, so a password is not required.

```
6. [student@workstation ~]$ ssh podsvc@serverb
```

7. ...output omitted...

```
[podsvc@serverb ~]$
```

8. Log in to the container registry using the `podman login` command.

```
9. [podsvc@serverb ~]$ podman login registry.lab.example.com
```



```
10. Username: admin
```

```
11. Password: redhat321
```

```
Login Succeeded!
```

12. Use the `skopeo inspect` command to view information about the `registry.lab.example.com/rhel8/mariadb-103` image. The following `skopeo inspect` command is very long and should be entered as a single line.

```
13. [podsvc@serverb ~]$ skopeo inspect
    docker://registry.lab.example.com/rhel8/mariadb-103
14. {
15.     "Name": "registry.lab.example.com/rhel8/mariadb-103",
16.     "Digest": "sha256:a95b...4816",
17.     "RepoTags": [
18.         "1-86",
19.         "1-102",
20.         "latest"
21.     ],
```

```
...output omitted...
```

The tag with the lowest number is 1-86.

Hide Solution

4. On `serverb`, as the `podsvc` user, create the `/home/podsvc/db_data` directory. Prepare the directory so that containers have read/write access. You will use this directory for persistent storage.

1. Create the `/home/podsvc/db_data` directory.

```
2. [podsvc@serverb ~]$ mkdir /home/podsvc/db_data
```

```
[podsvc@serverb ~]$
```

3. Set the access mode of the directory to `777` so that everyone has read/write access.

```
4. [podsvc@serverb ~]$ chmod 777 /home/podsvc/db_data
```

```
[podsvc@serverb ~]$
```

5. Hide Solution

6. On serverb, as the podsvc user, create a detached MariaDB container named inventorydb. Use the rhel8/mariadb-103 image from the registry.lab.example.com registry, specifying the tag with the lowest version number on that image, which you found in a preceding step. Map port 3306 in the container to port 13306 on the host. Mount the /home/podsvc/db_data directory on the host as /var/lib/mysql/data in the container. Declare the following variable values:

Variable	Value
MYSQL_USER	operator1
MYSQL_PASSWORD	redhat
MYSQL_DATABASE	inventory
MYSQL_ROOT_PASSWORD	redhat

7. You can copy and paste these parameters from the /home/podsvc/containers-review/variables file on serverb.
8. To confirm that the MariaDB database is running, use the mysql command. You can find this command in the /home/podsvc/containers-review/testdb.sh script. You can also directly run the script to test the database.
1. Use the podman run command to create the container. The following podman run command is very long and should be entered as a single line.

```
[podsvc@serverb ~]$ podman run -d --name inventorydb -p 13306:3306 -v /home/podsvc/db_data:/var/lib/mysql/data:Z -e MYSQL_USER=operator1 -e MYSQL_PASSWORD=redhat -e MYSQL_DATABASE=inventory -e MYSQL_ROOT_PASSWORD=redhat registry.lab.example.com/rhel8/mariadb-103:1-86
```

...output omitted...

3. Confirm that the database is running.

```
[podsvc@serverb ~]$ ~/containers-review/testdb.sh
```

5. Testing the access to the database...

SUCCESS

9. Hide Solution

10. On serverb, as the podsvc user, configure systemd so that the inventorydb container starts automatically with the server.

1. If you used `sudo` or `su` to log in as the `podsvc` user, then exit `serverb` and use the `ssh` command to directly log in to `serverb` as the `podsvc` user. Remember, `systemd` requires that the user open a direct session from the console or through SSH.

```
2. [student@workstation ~]$ ssh podsvc@serverb
```

```
3. ...output omitted...
```

```
[podsvc@serverb ~]$
```

4. Create the `~/.config/systemd/user/` directory.

```
5. [podsvc@serverb ~]$ mkdir -p ~/.config/systemd/user/
```

```
[podsvc@serverb ~]$
```

6. Use the `podman generate systemd` command to create the `systemd` unit file from the running container.

```
7. [podsvc@serverb ~]$ cd ~/.config/systemd/user/
```

```
8. [podsvc@serverb user]$ podman generate systemd --name inventorydb  
--files --new
```

```
/home/podsvc/.config/systemd/user/container-inventorydb.service
```

9. Stop and then delete the `inventorydb` container.

```
10. [podsvc@serverb user]$ podman stop inventorydb
```

```
11. 0d28f0e0a4118ff019691e34afe09b4d28ee526079b58d19f03b324bd04fd545
```

```
12. [podsvc@serverb user]$ podman rm inventorydb
```

```
0d28f0e0a4118ff019691e34afe09b4d28ee526079b58d19f03b324bd04fd545
```

13. Instruct `systemd` to reload its configuration, and then enable and start the `container-inventorydb` service.

```
14. [podsvc@serverb user]$ systemctl --user daemon-reload
```

```
15. [podsvc@serverb user]$ systemctl --user enable --now  
container-inventorydb.service
```

```
16. Created symlink  
/home/podsvc/.config/systemd/user/multi-user.target.wants/container-i  
nventorydb.service →  
/home/podsvc/.config/systemd/user/container-inventorydb.service.
```

```
Created symlink  
/home/podsvc/.config/systemd/user/default.target.wants/container-inve  
ntorydb.service →  
/home/podsvc/.config/systemd/user/container-inventorydb.service.
```

17. Confirm that the container is running.

```
18. [podsvc@serverb user]$ ~/containers-review/testdb.sh
```

19. Testing the access to the database...

20. SUCCESS

```
21. [podsvc@serverb user]$ podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
3ab24e7f000d	registry.lab.example.com/rhel8/mariadb-103:1-86	run-mysqld	47 seconds ago	Up 46 seconds ago	0.0.0.0:13306->3306/tcp inventorydb

```
3ab24e7f000d registry.lab.example.com/rhel8/mariadb-103:1-86
run-mysqld 47 seconds ago Up 46 seconds ago
0.0.0.0:13306->3306/tcp inventorydb
```

23. Run the `loginctl enable-linger` command for the user services to start automatically when the server starts.

```
24. [podsvc@serverb ~]$ loginctl enable-linger
```

```
[podsvc@serverb ~]$
```

25. Exit from serverb.

```
26. [podsvc@serverb ~]$ exit
```

27. logout

28. Connection to serverb closed.

```
[student@workstation ~]$
```

Lab 14: Configuring and Managing Server Security

1. From workstation, open an SSH session to serverb as student.

```
1. [student@workstation ~]$ ssh student@serverb
```

2. ...output omitted...

Hide Solution

2. Generate SSH keys for the student user on serverb using the `ssh-keygen` command. Do not protect the private key with a passphrase.

```
1. [student@serverb ~]$ ssh-keygen
```

2. Generating public/private rsa key pair.

```
3. Enter file in which to save the key (/home/student/.ssh/id_rsa):
Enter
```

4. Created directory '/home/student/.ssh'.

```
5. Enter passphrase (empty for no passphrase): Enter
```

```

6. Enter same passphrase again: Enter
7. Your identification has been saved in /home/student/.ssh/id_rsa.
8. Your public key has been saved in /home/student/.ssh/id_rsa.pub.
9. The key fingerprint is:
10. SHA256:1TPZ4TXYwiGWfExUGtRTHgFKQbF9hVuLa+VmH4vgkFY
    student@serverb.lab.example.com
11. The key's randomart image is:
12. +---[RSA 2048]---+
13. |                .+@BO**|
14. |                .=#+B*|
15. |                . X.*o=|
16. |                . E +.+ |
17. |                S o  +  |
18. |                + . o = |
19. |                . o o + +|
20. |                . . ..|
21. |                |
22. +----[SHA256]-----+

```

Hide Solution

3. On servera, configure the student user to accept login authentication using the SSH key pair you created for student on serverb. The student user on serverb should be able to log in to servera using SSH without entering a password. Use student as the password of the student user, when required.
 1. Use the ssh-copy-id command to transfer the public key of the SSH key pair of student on serverb to student on servera. Use student as the password of the student user, if prompted.

```

2. [student@serverb ~]$ ssh-copy-id student@servera
3. /usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed:
   "/home/student/.ssh/id_rsa.pub"
4. The authenticity of host 'servera (172.25.250.10)' can't be
   established.
5. ECDSA key fingerprint is
   SHA256:g/fIMtVzDWTbTi1l00wC30sL6cHmro9Tf563NxmeyyE.
6. Are you sure you want to continue connecting (yes/no)? yes
7. /usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s),
   to filter out any that are already installed
8. /usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you
   are prompted now it is to install the new keys
9. student@servera's password: student

```

10.
11. Number of key(s) added: 1
12.
13. Now try logging into the machine, with: "ssh 'student@servera'"

and check to make sure that only the key(s) you wanted were added.

14. Use the ssh command to verify that the student user can log in to servera from serverb without entering a password.

15. [student@serverb ~]\$ ssh student@servera
16. ...output omitted...

[student@servera ~]\$

4. Hide Solution

5. On servera, change the default SELinux mode to permissive.

1. Edit /etc/sysconfig/selinux to set the value of the parameter SELINUX to permissive. You can use the sudo vi /etc/sysconfig/selinux command to edit the configuration file as the superuser. Use the password student, if prompted.

2. ...output omitted...
3. #SELINUX=enforcing
4. SELINUX=permissive

...output omitted...

5. Use the sudo systemctl reboot command to reboot the system as the superuser.

6. [student@servera ~]\$ sudo systemctl reboot
7. Connection to servera closed by remote host.
8. Connection to servera closed.

[student@serverb ~]\$

6. Hide Solution

7. Configure serverb to automatically mount the home directory of the production5 user when the user logs in, using the network file system /home-directories/production5. This network file system is exported from servera.lab.example.com. Adjust the appropriate SELinux Boolean so that production5 can use the NFS-mounted home directory on serverb after

authenticating via SSH key-based authentication. The production5 user's password is redhat.

1. On serverb, use the `sudo -i` command to switch to the root user account.

```
[student@serverb ~]$ sudo -i
[sudo] password for student: student
```

```
[root@serverb ~]#
```

4. Install the autofs package.

```
[root@serverb ~]# yum install autofs
...output omitted...
7. Is this ok [y/N]: y
...output omitted...
9. Installed:
10.  autofs-1:5.1.4-29.el8.x86_64
11.
```

```
Complete!
```

12. Create the autofs master map file called `/etc/auto.master.d/production5.autofs` with the following content.

```
/- /etc/auto.production5
```

13. Retrieve the details of the production5 user to get the home directory path.

```
[root@serverb ~]# getent passwd production5
```

```
production5:x:5001:5001::/localhome/production5:/bin/bash
```

15. Create the `/etc/auto.production5` file with the following content.

```
/localhome/production5 -rw
servera.lab.example.com:/home-directories/production5
```

16. Restart the autofs service.

```
[root@serverb ~]# systemctl restart autofs
```

8. Hide Solution
9. On servera, verify that the production5 user is not able to log in to serverb using SSH public-key authentication. An SELinux Boolean causes this issue which you will fix in the following steps.

1. From workstation, open an SSH session to servera as student.

```
2. [student@workstation ~]$ ssh student@servera
3. ...output omitted...
```

```
[student@servera ~]$
```

4. Switch to the production5 user using the password redhat.

```
5. [student@servera ~]$ su - production5
6. Password: redhat
```

```
[production5@servera ~]$
```

7. Use the ssh-keygen command to generate the SSH keys as production5.

```
8. [production5@servera ~]$ ssh-keygen
9. Generating public/private rsa key pair.
10. Enter file in which to save the key (/home/production5/.ssh/id_rsa):
    Enter
11. Created directory '/home/production5/.ssh'.
12. Enter passphrase (empty for no passphrase): Enter
13. Enter same passphrase again: Enter
14. Your identification has been saved in /home/production5/.ssh/id_rsa.
15. Your public key has been saved in /home/production5/.ssh/id_rsa.pub.
16. The key fingerprint is:
17. SHA256:zmin1nmCt4H8LA+4FPimtdg81n17ATbInUFW3HSPxk4
    production5@servera.lab.example.com
18. The key's randomart image is:
19. +---[RSA 2048]----+
20. |      .oo.o. .  |
21. |      .. . .o o |
22. |      . o o   E .|
23. |      . o *    + |
24. |      . . .So   . |
25. |      . + =    . |
26. |      *. * += .  |
```



```
27. | 0o+***.o |
28. | o.=o.=** |
```

```
+-----[SHA256]-----+
```

29. Use the `ssh-copy-id` command to transfer the public key of the SSH key pair of `production5` on `servera` to `production5` on `serverb`. Use `redhat` as the password of the `production5` user, if prompted.

```
30. [production5@servera ~]$ ssh-copy-id production5@serverb
31. /usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed:
    "/home/production5/.ssh/id_rsa.pub"
32. The authenticity of host 'serverb (172.25.250.11)' can't be
    established.
33. ECDSA key fingerprint is
    SHA256:ciCkaRWF4g6eR9nSdPxQ7KL8czpViXal6BousK544TY.
34. Are you sure you want to continue connecting (yes/no)? yes
35. /usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s),
    to filter out any that are already installed
36. /usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you
    are prompted now it is to install the new keys
37. production5@serverb's password: redhat
38.
39. Number of key(s) added: 1
40.
41. Now try logging into the machine, with:  "ssh 'production5@serverb'"
```

and check to make sure that only the key(s) you wanted were added.

42. Use the SSH public key-based authentication instead of password-based authentication to log in to `serverb` as `production5`. This command should fail.

```
43. [production5@servera ~]$ ssh -o pubkeyauthentication=yes \
44. -o passwordauthentication=no production5@serverb
```

```
production5@serverb: Permission denied
(publickey,gssapi-keyex,gssapi-with-mic,password).
```

10. Hide Solution

11. Set the appropriate SELinux Boolean setting on `serverb`, so that `production5` can log in to `serverb` using the SSH public key-based authentication and use the home directory.

1. On serverb as root, set the use_nfs_home_dirs SELinux Boolean to true.

```
[root@serverb ~]# setsebool -P use_nfs_home_dirs true
```

2. Use the SSH public key-based authentication instead of password-based authentication to log in to serverb as production5. This command should succeed.

3. [production5@servera ~]\$ ssh -o pubkeyauthentication=yes \
4. -o passwordauthentication=no production5@serverb
5. ...output omitted...

```
[production5@serverb ~]$
```

12.Hide Solution

- 13.On serverb, adjust the firewall settings so that SSH connections originating from servera are rejected. The servera system uses the IPv4 address 172.25.250.10.

1. Use the firewall-cmd command to add the IPv4 address of servera to the firewalld zone called block.

2. [root@serverb ~]# firewall-cmd --add-source=172.25.250.10/32 \
3. --zone=block --permanent

```
success
```

4. Use the firewall-cmd --reload command to reload the changes in the firewall settings.

5. [root@serverb ~]# firewall-cmd --reload

```
success
```

14.Hide Solution

- 15.On serverb, investigate and fix the issue with the Apache HTTPD daemon, which is configured to listen on port 30080/TCP, but which fails to start. Adjust the firewall settings appropriately so that port 30080/TCP is open for incoming connections.

1. Use the systemctl command to restart the httpd service. This command fails to restart the service.

2. [root@serverb ~]# systemctl restart httpd.service
3. Job for httpd.service failed because the control process exited with error code.

See "systemctl status httpd.service" and "journalctl -xe" for details.

4. Use the `systemctl status` command to investigate the reason for the failure of the `httpd` service.

```
5. [root@serverb ~]# systemctl status httpd.service
6. • httpd.service - The Apache HTTP Server
7.    Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled;
            vendor preset: disabled)
8.    Active: failed (Result: exit-code) since Mon 2019-04-15 06:42:41
            EDT; 5min ago
9.    Docs: man:httpd.service(8)
10.   Process: 27313 ExecStart=/usr/sbin/httpd $OPTIONS -DFOREGROUND
            (code=exited, status=1/FAILURE)
11.  Main PID: 27313 (code=exited, status=1/FAILURE)
12.   Status: "Reading configuration..."
13.
14. Apr 15 06:42:41 serverb.lab.example.com systemd[1]: Starting The
            Apache HTTP Server...
15. Apr 15 06:42:41 serverb.lab.example.com httpd[27313]: (13)Permission
            denied: AH00072: make_sock: could not bind to address [::]:30080
16. Apr 15 06:42:41 serverb.lab.example.com httpd[27313]: (13)Permission
            denied: AH00072: make_sock: could not bind to address 0.0.0.0:30080
17. Apr 15 06:42:41 serverb.lab.example.com httpd[27313]: no listening
            sockets available, shutting down
18. Apr 15 06:42:41 serverb.lab.example.com httpd[27313]: AH00015: Unable
            to open logs
19. Apr 15 06:42:41 serverb.lab.example.com systemd[1]: httpd.service:
            Main process exited, code=exited, status=1/FAILURE
20. Apr 15 06:42:41 serverb.lab.example.com systemd[1]: httpd.service:
            Failed with result 'exit-code'.
```

```
Apr 15 06:42:41 serverb.lab.example.com systemd[1]: Failed to start
The Apache HTTP Server.
```

Notice the permission error in the preceding output, which signifies that the `httpd` daemon failed to bind to port `30080/TCP`. The SELinux policy can be a potential restriction for an application to bind to a port. Press **q** to quit the preceding `systemctl` command.

21. Use the `sealert` command to determine if an SELinux policy is preventing `httpd` from binding to port `30080/TCP`.

```
22. [root@serverb ~]# sealert -a /var/log/audit/audit.log
```

```

23.100% done
24.found 1 alerts in /var/log/audit/audit.log
25.-----
26.
27.SELinux is preventing /usr/sbin/httpd from name_bind access on the
   tcp_socket port 30080.
28.
29.***** Plugin bind_ports (92.2 confidence) suggests
   *****
30.
31.If you want to allow /usr/sbin/httpd to bind to network port 30080
32.Then you need to modify the port type.
33.Do
34.# semanage port -a -t PORT_TYPE -p tcp 30080
35.   where PORT_TYPE is one of the following: http_cache_port_t,
   http_port_t, jboss_management_port_t, jboss_messaging_port_t,
   ntop_port_t, puppet_port_t.

```

...output omitted...

The preceding log message reveals that the port 30080/TCP does not have the appropriate SELinux context `http_port_t`, causing SELinux to prevent `httpd` to bind to this port. The log message also produces the syntax of the `semanage port` command so that you can easily fix the issue.

36. Use the `semanage port` command to set the appropriate SELinux context on the port 30080/TCP for `httpd` to bind to it.

```
[root@serverb ~]# semanage port -a -t http_port_t -p tcp 30080
```

37. Use the `systemctl` command to restart `httpd`. This command should successfully restart the service.

```
[root@serverb ~]# systemctl restart httpd
```

38. Add the port 30080/TCP to the default `firewalld` zone called `public`.

```

39.[root@serverb ~]# firewall-cmd --add-port=30080/tcp --permanent
40.success
41.[root@serverb ~]# firewall-cmd --reload

```

success

42. Exit the root user's shell.

43. [root@serverb ~]# exit

logout

44. Log off from serverb.

45. [student@serverb ~]\$ exit

46. logout

Connection to serverb closed.